

## Clase 8

### 1. Crear la siguientes tablas externas en la base de datos f1 en hive:

- driver\_results (driver\_forename, driver\_surname, driver\_nationality, points)
- constructor\_results (constructorRef, cons\_name, cons\_nationality, url, points)

CREATE DATABASE f1;

```
hive> create database f1;
OK
Time taken: 1.217 seconds
hive> show databases;
OK
default
f1
tripdata
Time taken: 0.042 seconds, Fetched: 3 row(s)
hive>
```

```
CREATE EXTERNAL TABLE f1.driver_results(driver_forename string, driver_surname
string, driver_nationality string, points float)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/tables/external/f1/driver_results';
```

```
CREATE EXTERNAL TABLE f1.constructor_results(constructorRef string, cons_name
string, cons_nationality string, url string, points float)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/tables/external/f1/constructor_results';
```

```
hive> show tables;
OK
constructor_results
driver_results
Time taken: 0.057 seconds, Fetched: 2 row(s)
hive>
```

### 2. En Hive, mostrar el esquema de driver\_results y constructor\_results

```
hive> describe driver_results;
OK
driver_forename      string
driver_surname       string
driver_nationality   string
points               float
Time taken: 0.064 seconds, Fetched: 4 row(s)
hive>
```

```
hive> describe constructor_results;
OK
constructorref       string
cons_name            string
cons_nationality     string
url                  string
points               float
Time taken: 0.076 seconds, Fetched: 5 row(s)
hive>
```

### 3. Crear un archivo .bash que permita descargar los archivos mencionados abajo e ingestarlos en HDFS:

results.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/results.csv>

drivers.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/drivers.csv>

constructors.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/constructors.csv>

racers.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/races.csv>

```
hadoop@d41c15beb563:~/scripts$ cat ingest_clase_ocho.sh
ruta='/home/hadoop/landing/'
ruta_hdfs='/home/hadoop/hadoop/bin/'

rm -f "${ruta}”*.*

wget -P "${ruta}" https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/results.csv
wget -P "${ruta}" https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/drivers.csv
wget -P "${ruta}" https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/constructors.csv
wget -P "${ruta}" https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/races.csv

"${ruta_hdfs}hdfs" dfs -rm /ingest/*. *

"${ruta_hdfs}hdfs" dfs -put "${ruta}”*.* /ingest
hadoop@d41c15beb563:~/scripts$
```

### 4. Generar un archivo .py que permita, mediante Spark:

- insertar en la tabla driver\_results los corredores con mayor cantidad de puntos en la historia.
- insertar en la tabla constructor\_result quienes obtuvieron más puntos en el Spanish Grand Prix en el año 1991.

```
hadoop@ec27dbd59e9:~/scripts$ cat transform_clase_ocho.py
from pyspark.context import SparkContext
from pyspark.sql.session import SparkSession
from pyspark.sql import HiveContext
sc = SparkContext('local')
spark = SparkSession(sc)
hc = HiveContext(sc)

#leo csv de HDFS y lo cargo en un dataframe
df_drivers = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/drivers.csv")
df_results = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/results.csv")
df_constructors = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/constructors.csv")
df_races = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/races.csv")

#creamos una vista de los DF
df_drivers.createOrReplaceTempView("drivers")
df_results.createOrReplaceTempView("results")
df_constructors.createOrReplaceTempView("constructors")
df_races.createOrReplaceTempView("races")

#unimos y filtramos el DF para corredores con mayor cantidad de puntos en la historia
df_drivers2 = spark.sql("select cast(driverId as int) as id, cast(forename as string) as driver_forename, cast(surname as string) as driver_surname, cast(nationality as string) as driver_nationality from drivers")
df_results2 = spark.sql("select cast(driverId as int) as id, cast(points as float) from results")

df_drivers2.createOrReplaceTempView("drivers_filtrados")
df_results2.createOrReplaceTempView("results_filtrados")

df_final = spark.sql("select drivers_filtrados.driver_forename, drivers_filtrados.driver_surname, drivers_filtrados.driver_nationality, SUM(results_filtrados.points) as points from drivers_filtrados inner join results_filtrados on results_filtrados.id where points > 0 group by drivers_filtrados.driver_forename, drivers_filtrados.driver_surname, drivers_filtrados.driver_nationality order by points desc limit 10")

df_final.createOrReplaceTempView("corredores_puntos")

#insertamos el DF en la tabla f1.driver_results
spark.sql("insert into f1.driver_results select * from corredores_puntos")

#unimos y filtramos el DF para quienes obtuvieron más puntos en el Spanish Grand Prix en el año 1991
df_constructors2 = spark.sql("select cast(constructorId as string), cast(constructorRef as string), cast(name as string) as cons_name, cast(nationality as string) as cons_nationality, cast(url as string) from constructors")
df_races2 = spark.sql("select cast(raceId as int), cast(raceId as int), cast(points as float) from results")
df_races2 = spark.sql("select cast(raceId as int), cast(year as int), cast(name as string) from races")

df_constructors2.createOrReplaceTempView("constructors_filtrados")
df_results2.createOrReplaceTempView("results_filtrados")
df_races2.createOrReplaceTempView("races_filtrados")

df_results_races = spark.sql("select * from results_filtrados inner join races_filtrados on results_filtrados.raceId = races_filtrados.raceId inner join constructors_filtrados on constructors_filtrados.constructorId = results_filtrados.constructorId")
df_results_races.createOrReplaceTempView("join_tablas")

df_final = spark.sql("select constructorRef, cons_name, cons_nationality, url, SUM(points) as points from join_tablas where year = 1991 and name = 'Spanish Grand Prix' and points > 0 group by constructorRef, cons_name, cons_nationality, url")
df_final.createOrReplaceTempView("constructores_puntos")

#insertamos el DF en la tabla f1.constructor_results
spark.sql("insert into f1.constructor_results select * from constructores_puntos")
hadoop@ec27dbd59e9:~/scripts$
```

5. Realizar un proceso automático en Airflow que orqueste los archivos creados en los puntos 3 y 4. Correrlo y mostrar una captura de pantalla (del DAG y del resultado en la base de datos)

```
1 from datetime import timedelta
2 from airflow import DAG
3 from airflow.operators.bash import BashOperator
4 from airflow.operators.dummy import DummyOperator
5 from airflow.utils.dates import days_ago
6
7 args = {
8     'owner': 'airflow',
9 }
10
11 with DAG(
12     dag_id='ingest-transform-load-clase-ocho',
13     default_args=args,
14     schedule_interval='0 0 * * *',
15     start_date=days_ago(2),
16     dagrun_timeout=timedelta(minutes=60),
17     tags=['ingest', 'transform'],
18     params={"example_key": "example_value"},
19 ) as dag:
20
21     inicia_proceso = DummyOperator(
22         task_id='inicia_proceso',
23     )
24
25     finaliza_proceso = DummyOperator(
26         task_id='finaliza_proceso',
27     )
28
29     ingest = BashOperator(
30         task_id='ingest',
31         bash_command='/usr/bin/sh /home/hadoop/scripts/ingest_clase_ocho.sh ',
32     )
33
34
35     transform_load = BashOperator(
36         task_id='transform_load',
37         bash_command='ssh hadoop@172.17.0.2 /home/hadoop/spark/bin/spark-submit --files /home/hadoop/hive/conf/hive-site.xml /home/hadoop/scripts/transform_clase_ocho.py ',
38     )
39
40
41
42     inicia_proceso >> ingest >> transform_load >> finaliza_proceso
```

Airflow

DAAGs

Security

Browse

Admin

Docs

02:52 UTC

AA

DAG: ingest-transform-load-clase-ocho

queuedSchedule: 0 0 \* \* \*Next Run: 2024-05-23, 00:00:00

Grid

Graph

Calendar

Task Duration

Task Tries

Landing Times

Gantt

Details

<> Code

Audit Log

2024-05-23T02:50:26Z

Runs

25

Run

manual\_\_2024-05-23T02:50:25.167090+00:00

Layout

Left > Right

Update

Find Task...

BashOperator

DummyOperator

queued

running

success

failed

up\_for\_retry

up\_for\_reschedule

upstream\_failed

skipped

scheduled

deferred

no\_status

Auto-refresh

inicia\_proceso

→

ingest

→

transform\_load

→

finaliza\_proceso

select \* from f1\_driver\_results dt

Enter a SQL expression to filter results (use Ctrl+Space)

	ABC driver_forename	ABC driver_surname	ABC driver_nationality	123 points	
1	Lewis	Hamilton	British	4,308.5	
2	Sebastian	Vettel	German	3,077	
3	Fernando	Alonso	Spanish	2,021	
4	Kimi	Räikkönen	Finnish	1,873	
5	Max	Verstappen	Dutch	1,792.5	
6	Valtteri	Bottas	Finnish	1,775	
7	Nico	Rosberg	German	1,594.5	
8	Michael	Schumacher	German	1,566	
9	Daniel	Ricciardo	Australian	1,289	
10	Jenson	Button	British	1,235	

SQL

select \* from f1\_constructor\_results order by points desc

Enter a SQL expression to filter results (use Ctrl+Space)

	ABC constructorref	ABC cons_name	ABC cons_nationality	ABC url	123 points
1	williams	Williams	British	<a href="http://en.wikipedia.org/wiki/Williams_Grand_Prix_Engineering">http://en.wikipedia.org/wiki/Williams_Grand_Prix_Engineering</a>	14
2	ferrari	Ferrari	Italian	<a href="http://en.wikipedia.org/wiki/Scuderia_Ferrari">http://en.wikipedia.org/wiki/Scuderia_Ferrari</a>	9
3	mclaren	McLaren	British	<a href="http://en.wikipedia.org/wiki/McLaren">http://en.wikipedia.org/wiki/McLaren</a>	2
4	benetton	Benetton	Italian	<a href="http://en.wikipedia.org/wiki/Benetton_Formula">http://en.wikipedia.org/wiki/Benetton_Formula</a>	1